

Module 1: Introduction to Unsupervised Learning.

1. **Introduction to Unsupervised Learning:** Unsupervised learning is a type of machine learning that looks for previously undetected patterns in a dataset with no pre-existing labels and with a minimum of human supervision. For example, suppose you have a collection of articles from various news sources. With unsupervised learning, you could cluster these documents into groups based on their content, without knowing in advance what these groups might be. This could be useful for a news agency wanting to categorize news stories automatically.
2. **Common Approaches to Clustering:** Clustering is a technique used in unsupervised learning that involves grouping data points. In essence, this technique involves identifying a subset of the data points in the dataset that are like each other. For instance, a marketing executive might want to segment the company's customer base into distinct groups with similar buying behaviors. They could then tailor the marketing strategy to each of these groups to increase sales.
3. **Some notes on the demonstrations:** This would generally involve explaining the practical applications and demonstrations of unsupervised learning algorithms. For example, a walkthrough of how to apply the K-Means clustering algorithm to a dataset in Python, demonstrating how the algorithm works step-by-step, and discussing how to interpret the results.

Module 2: K-Means Algorithm.

1. **K-Means - Getting Started:** K-means is a type of partitioning clustering that separates the data into K non-overlapping subsets (clusters). Each data point belongs to the cluster with the nearest mean (center). For instance, a business might use K-means to segment its customer base into different groups for targeted marketing.
2. **K-Means Algorithm:** The K-means algorithm assigns each point to the cluster whose center (also called centroid) is nearest. The center is the average of all the points in the cluster — that is, its coordinates are the arithmetic mean for each dimension separately over all the points in the cluster.
3. **Convergence & Performance Aspects:** K-means is guaranteed to converge, i.e., to find a local optimum. However, the solution might not be globally optimal, i.e., the best possible clustering. The quality of the final solution can depend on the initial set of clusters, and therefore, the algorithm is often run multiple times with different random initializations.
4. **K-Means Python Code:** This would cover the practical implementation of the K-means algorithm using Python, using libraries such as scikit-learn.
5. **Choosing Number (K) of Clusters:** One of the challenges in K-means is choosing the right number of clusters (K). One common approach is the Elbow method, which involves plotting the explained variation as a function of the number of clusters and picking the elbow of the curve as the number of clusters to use.

6. **K-Means Algorithm - Some Enhancements:** These might include methods to help choose K (such as the silhouette method), or ways to ensure a more globally-optimal solution (such as K-means++).
7. **K-Means Algorithm - Online Version (with sequential update):** This variant of K-means, often used in large-scale applications, updates the cluster centers one instance at a time rather than all at once.
8. **Mini-batch K-Means Algorithm:** This is a variant of the K-Means algorithm that uses mini-batches to reduce the computation time, while still attempting to optimise the same objective function.
9. **K-Means Algorithm & Outliers:** K-means can be sensitive to outliers, since they can pull the cluster center towards them. Techniques to mitigate this include using a variant of K-means that is more robust to outliers, or pre-processing the data to remove outliers.

Module 3: Expectation - Maximization Algorithm.

1. **Introduction:** Expectation-Maximization (EM) is a statistical algorithm for finding maximum likelihood estimates of parameters in probabilistic models, where the model depends on unseen latent variables.
2. **What is Mixtures of Gaussian?:** This refers to a probabilistic model for representing the presence of sub-populations within an overall population. For example, in a dataset of human heights, a mixture of two Gaussian distributions might be used to represent the heights of adult males and adult females.
3. **Clustering as finding ML parameters for MoG:** The EM algorithm can be used to determine the parameters (mean and standard deviation) for each Gaussian in a mixture, effectively clustering the data. In the human heights example, this would involve finding the mean and standard deviation of heights for adult males and adult females.
4. **Expectation Maximization Algorithm - Initial Formulation:** This refers to the initialization of the EM algorithm, which often involves setting initial values for the parameters of the Gaussian distributions, then iteratively updating these parameters to maximize the likelihood of the observed data.
5. **Expectation Maximization Algorithm - A General Formulation:** The EM algorithm consists of two steps: an expectation step, which calculates the expected value of the log likelihood function under the current estimate for the parameters, and a maximization step, which computes parameters maximizing the expected log-likelihood.
6. **EM Algorithm and K-Means Algorithm - Connection:** The K-means algorithm can be viewed as a variant of the EM algorithm where each data point is assigned exclusively to one cluster (hard assignment), whereas in the general EM algorithm, data points can belong to each cluster to a certain degree (soft assignment).
7. **Closing Points on EM Algorithm:** This would typically involve discussing the strengths and weaknesses of the EM algorithm, its applications, and how it compares to other clustering algorithms.

The use case for EM includes scenarios where data appears to be drawn from a mixture of different distributions. For example, in customer segmentation, different customer behaviors

may follow different probability distributions, and EM can be used to identify these different groups.

Module 4: Hierarchical Clustering Algorithms.

1. **Hierarchical Clustering Algorithms - Introduction:** Hierarchical clustering is a method of cluster analysis which seeks to build a hierarchy of clusters. The end result is a set of clusters that may be visualized in a dendrogram, a tree-like diagram that shows the sequence of merges or splits.
2. **What are Divisive & Agglomerative Clustering Algorithms?:** Hierarchical clustering can be Divisive (top-down) where all observations start in one cluster and splits are performed recursively as one moves down the hierarchy, or Agglomerative (bottom-up) where each observation starts in its own cluster, and pairs of clusters are merged as one moves up the hierarchy.
3. **Single Linkage Clustering:** This is a type of agglomerative hierarchical clustering that considers the shortest distance between clusters for merging. The distance between two clusters is defined as the shortest distance between two points in each cluster.
4. **Complete Linkage Clustering:** This is another type of agglomerative hierarchical clustering. Here, the distance between two clusters is defined as the longest distance between two points in each cluster.
5. **Mean and Average Linkage Clustering:** Mean linkage clustering uses the average of the distances of each point of one cluster to every point of the other cluster. Average linkage clustering is similar, but it uses the average of all pairwise distances between the elements in the two clusters.

A common use case for hierarchical clustering is in biological sciences, where it's used to group similar species of animals or plants. It's also used in social network analysis to find communities or groups of similar interests. These techniques allow for a multi-level clustering which can be very insightful and allows for variable cluster sizes.

Module 5: Density Based Clustering Algorithms.

1. **Introduction to Density Based Clustering:** Density-based clustering refers to a family of unsupervised machine learning algorithms that seek to cluster data points based on the idea of density. The basic idea is that a cluster is a high-density region surrounded by a region of lower density.
2. **DBSCAN - Understanding Terminologies - Part I:** DBSCAN (Density-Based Spatial Clustering of Applications with Noise) is a popular density-based clustering algorithm that separates high-density regions separated by regions of low density. Core concepts include the definition of a point's ϵ -neighborhood (all points within a certain radius), the minimum number of points required to form a cluster, and the concept of core, border, and noise points.
3. **DBSCAN - Understanding Terminologies - Part II:** This would delve deeper into the core concepts of DBSCAN, such as how clusters are formed and how noise points are treated.
4. **DBSCAN Algorithm:** The DBSCAN algorithm starts with an arbitrary point, checks if there are at least 'MinPts' nearby points (within distance 'EPS'), if so it starts a new

cluster, otherwise, the point is labeled as noise. The process is repeated until all points are processed.

5. **DBSCAN Algorithm - Some Points:** This likely discusses some important characteristics of DBSCAN, such as its ability to find arbitrarily shaped clusters, its need for parameter tuning, and its robustness to the initial order of the input data.

DBSCAN can be used in many fields where data can be understood as points in a space, such as geospatial analysis, anomaly detection, and image segmentation. For example, in geospatial analysis, DBSCAN can be used to find areas of high building density.